

Neural Network Acceleration on GPUs

By

Ibrahim Ahmed Lone
Muhammad Yaseen
Abdul Hadi

22i-1193
22i-1221
22i-1333

How much speedup is possible on a MNIST Neural Network using GPU—and at what accuracy cost?





Agenda

- Version 1: Sequential CPU
- Version 2: Naïve CUDA
- Version 3: Optimized CUDA
- Version 4: Tensor-Core Acceleration
- Comparative Results
- Conclusion

Version 1



Version 1: Sequential CPU

1. Results:

Total training time: 22.233 s

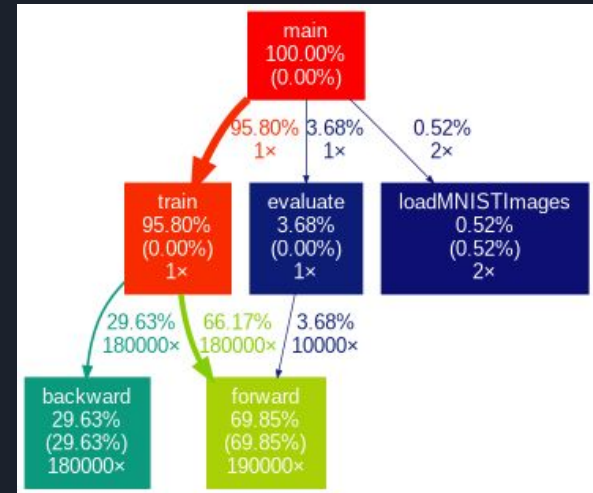
Test Accuracy: 96.78%.

End-to-end execution: 23.589 s.

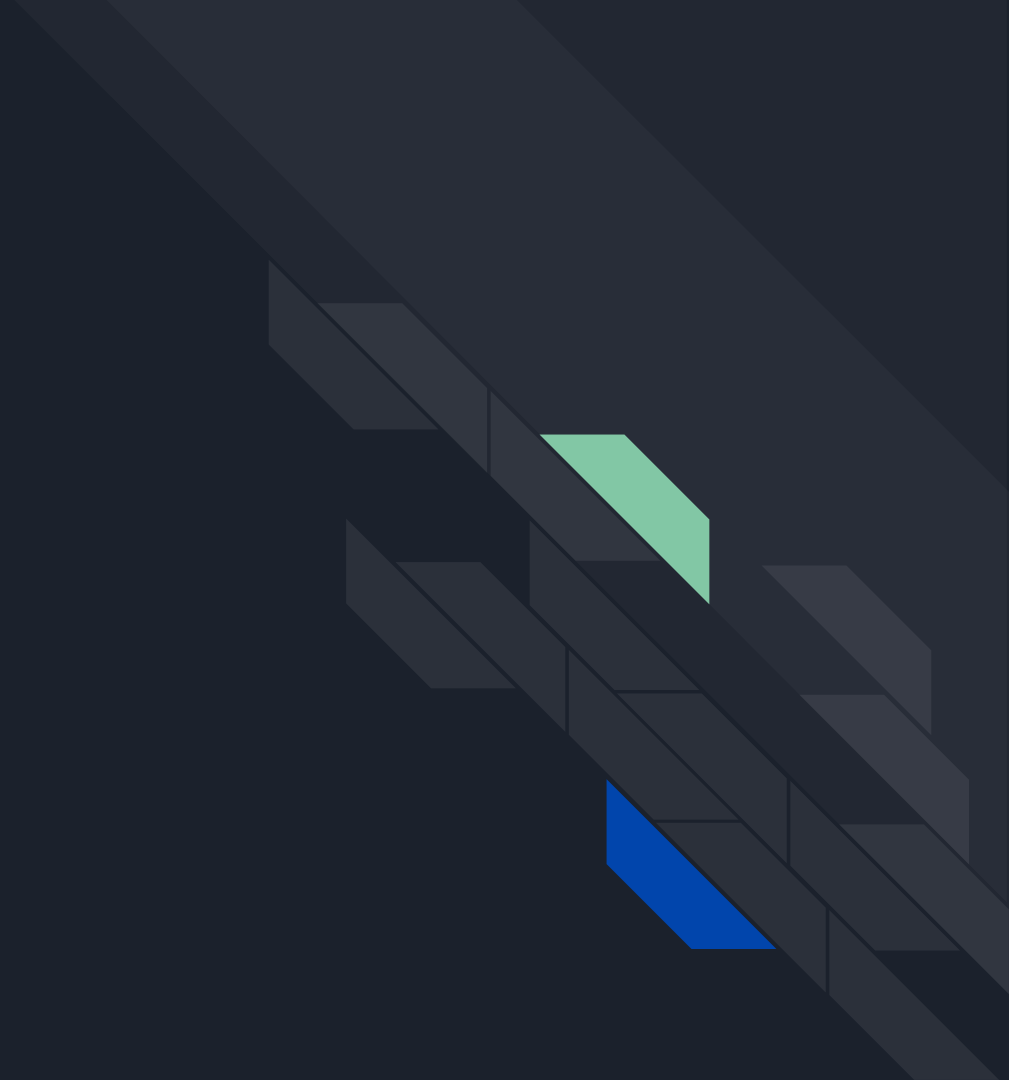
2. Observations:

Loop-carried Dependencies
(forward/backward).

Serves as the baseline for the speedup.



Version 2





Version 2: Naïve CUDA

Key changes:

- cudaMalloc/cudaMemcpy.

- Two kernels:

 - forward/backward pass.

Observations:

- GPU overhead > compute

- Low Occupancy

- Memory-bound operations are the bottleneck now

Results:

- Total training time: 206.037 s

- Test Accuracy: 96.58%

- Execution time: 207.567 s

Version 3





Version 3: Optimized CUDA

Key changes:

Launch & Occupancy: tuned dim3
grid/block (256–512 threads)

Communication: batched
H2D/D2H, streams overlap

Memory: shared-memory,
coalesced access

Results:

Total training time: 0.856s

Test Loss: 0.3440

Execution time: 1.097s

Observations:

~210 speedup in its training
compared with V2.

Memory-bound operations are
now well-optimized, with
compute-bound kernels being the
bottleneck now.

Version 4





Version 4: Tensor-Core Acceleration

Key changes:

Precision: FP16 inputs/weights,
FP32 accumulations
API: NVIDIA WMMA (16×16 warp
tiles)
Alignment: data padded to 16×16
blocks

Results:

Total training time: 0.262 s
Test accuracy: 88.76%
Execution time: 0.435 s

Observations:

~3.3× faster than V3; ~54× faster
than V1.

Occupancy still high; tensor cores
peak compute.



Comparative Results

Version	Train Time (s)	Exec Time (s)	Test Acc. (%)	Speedup vs. V1
V1	22.233	23.589	96.78	1×
V2	206.037	207.567	96.58	0.11×
V3	0.856	1.097	90.27	21.5×
V4	0.262	0.435	88.76	54.2×




Training Time and Execution Time





Accuracy and SpeedUp





Conclusion

Moving from CPU→naïve CUDA (V1→V2) without algorithmic changes delivers poor speedup due to overheads.

Advanced optimizations (shared memory, streams, WMMA) add significant code complexity and tuning effort - though yielding diminishing results.

The performance improvement trajectory showed a distinct drop from 23.6 s on CPU (V1) down to 0.44s using tensor cores (V4) at a speedup of $\sim 54\times$. The accuracy drops slightly when one is too aggressive on optimizations - speedup.